



UNIVERSITY OF
GOTHENBURG

SBX (local) IT infrastructure

Herbert Lange

October 22, 2024

SPRÅKBANKEN TEXT

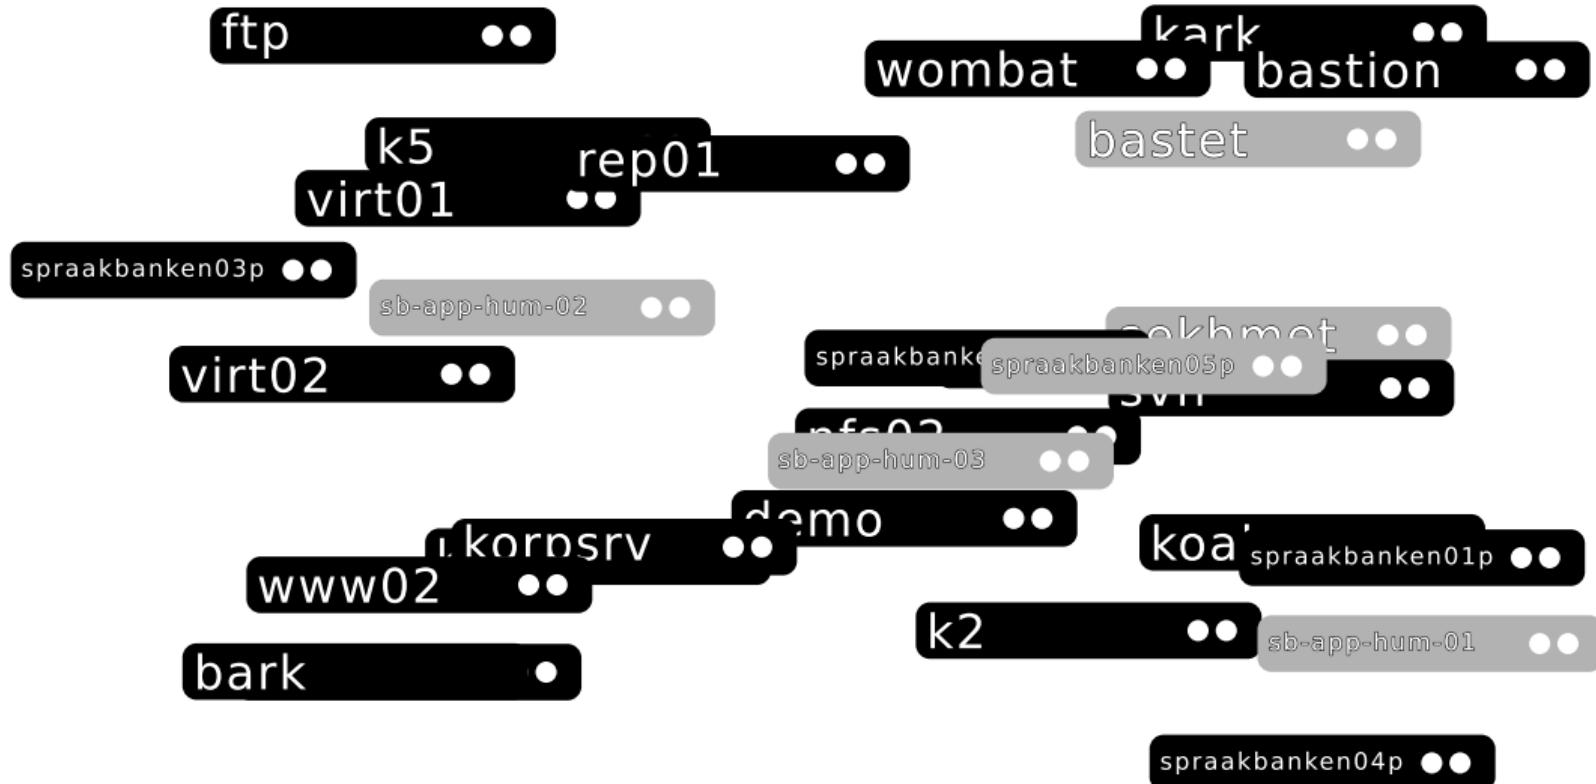
Quiz

Join on menti.com code **6729 7220**



Inventory

Servers



Servers – Clustered



Servers – End of life/Out of service

bastet



sekhemet



k5



sb-app-hum-01



sb-app-hum-02



sb-app-hum-03



Servers – Physical

bastion ••

virt01 ••

backup ••

nfs01 ••

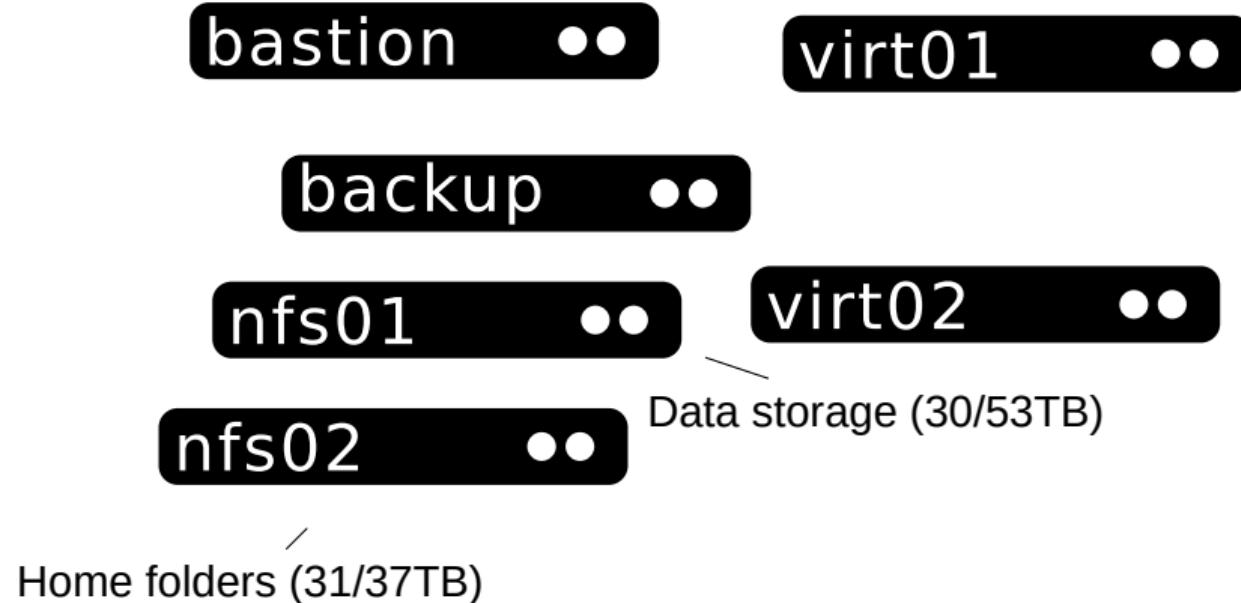
virt02 ••

nfs02 ••

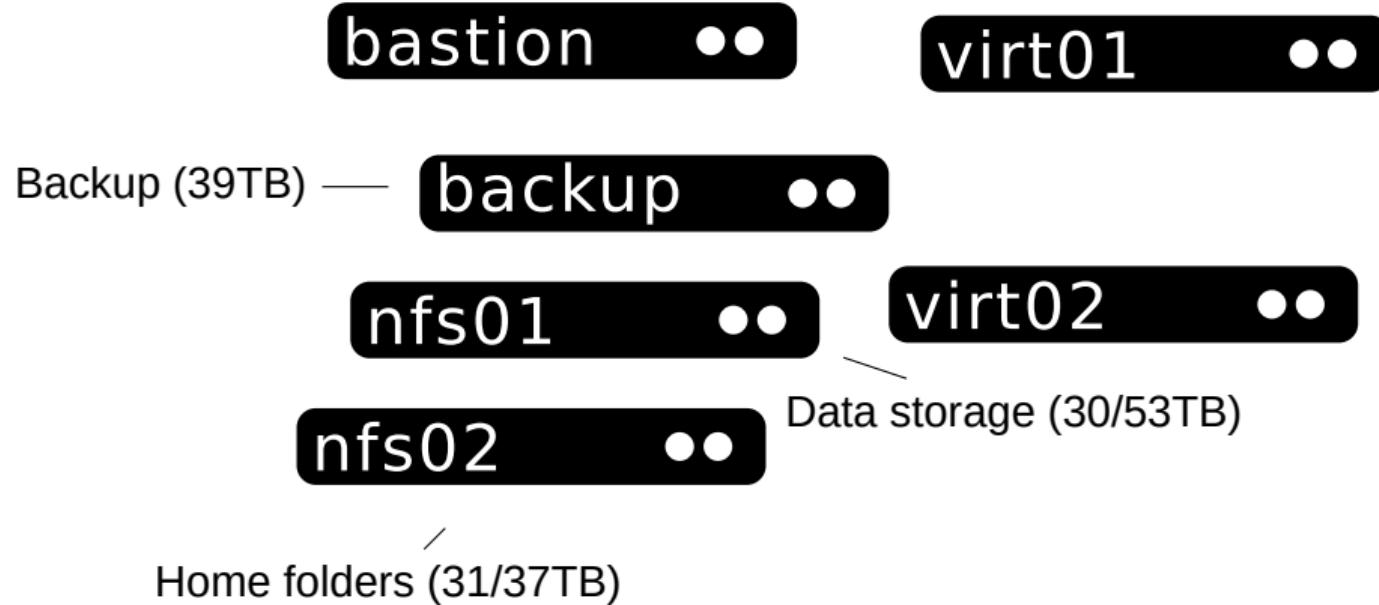
Servers – Physical



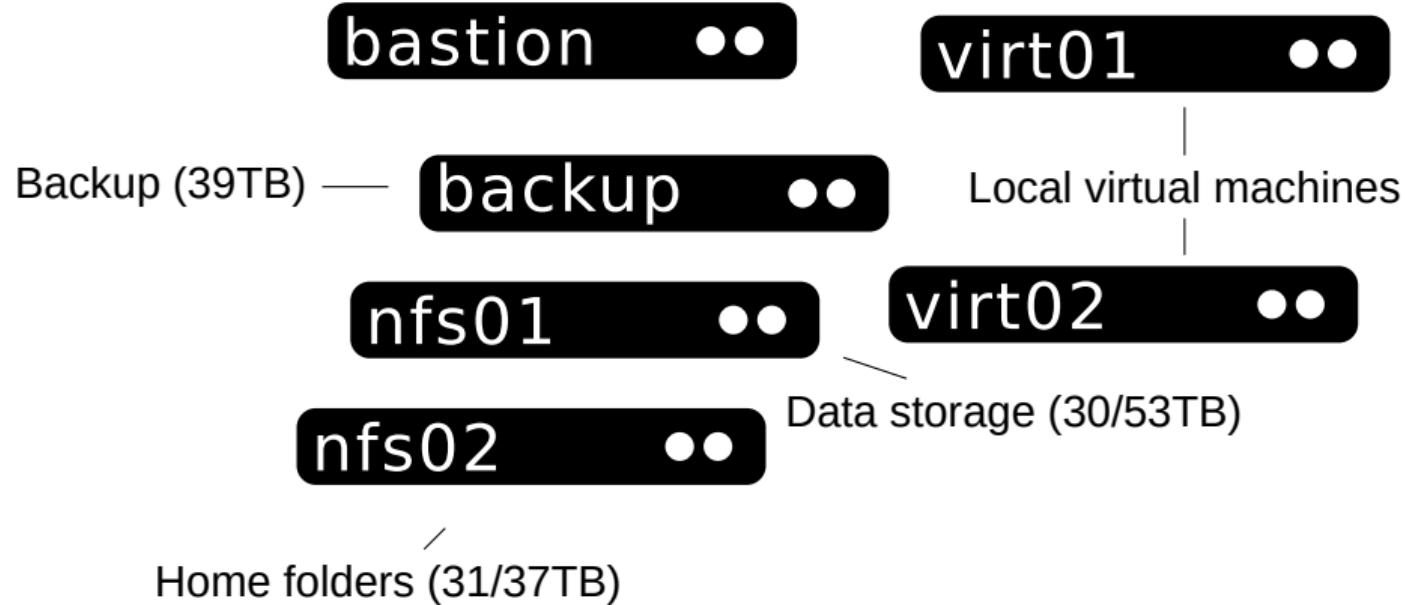
Servers – Physical



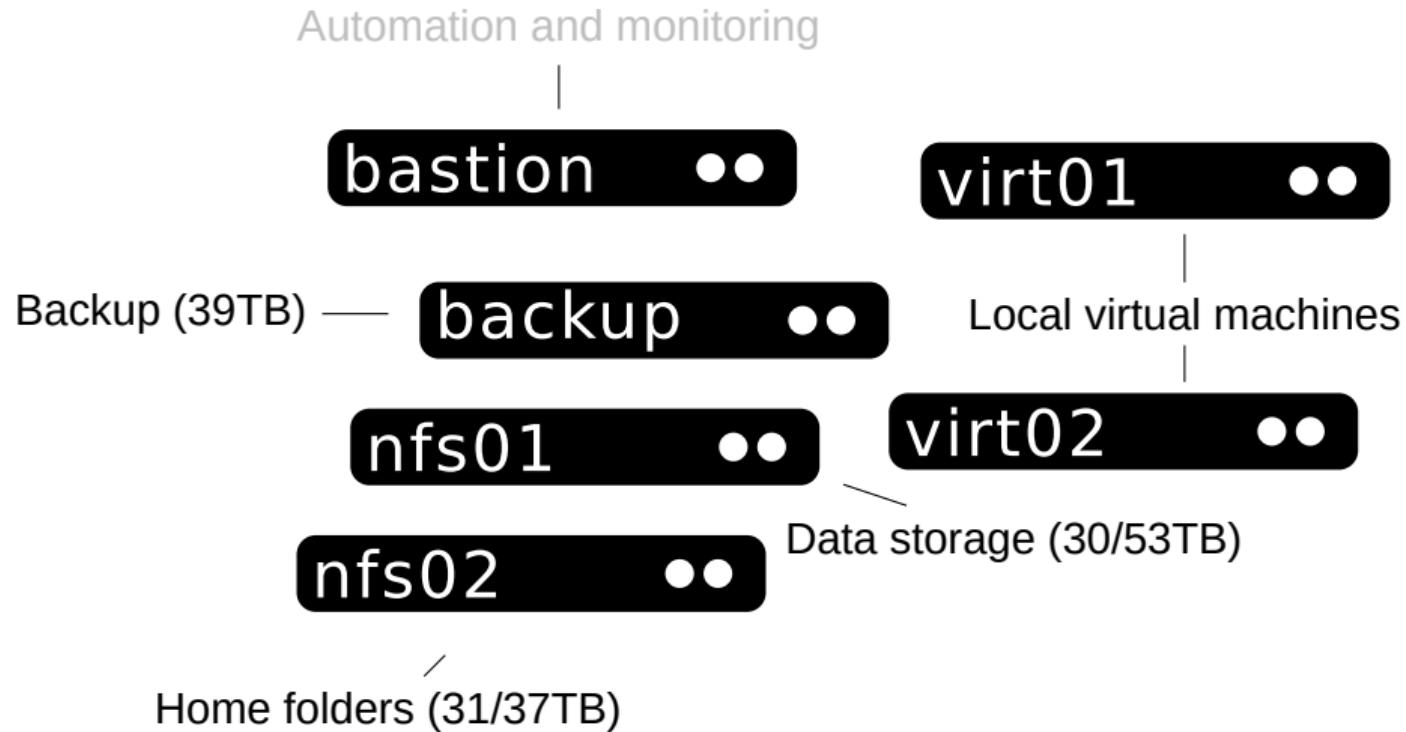
Servers – Physical



Servers – Physical



Servers – Physical



Servers – Web

rep01



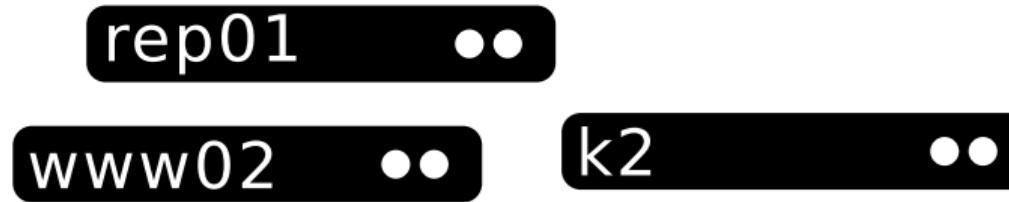
www02



k2

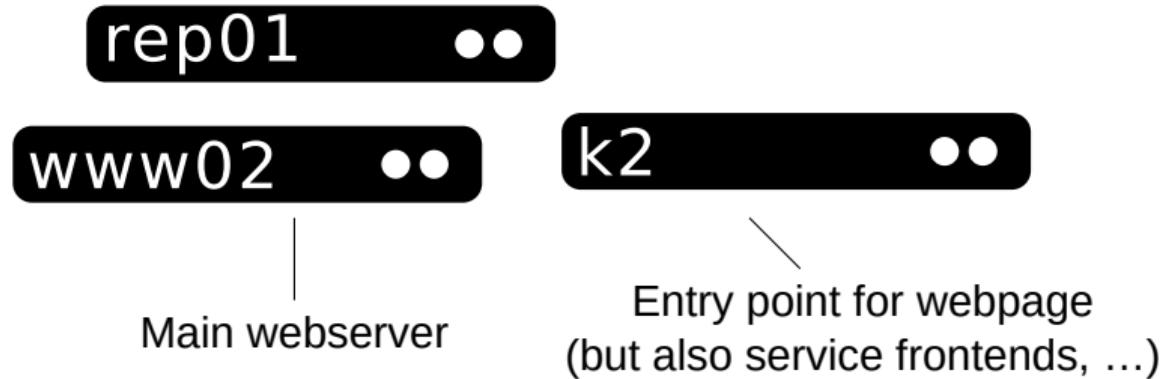


Servers – Web

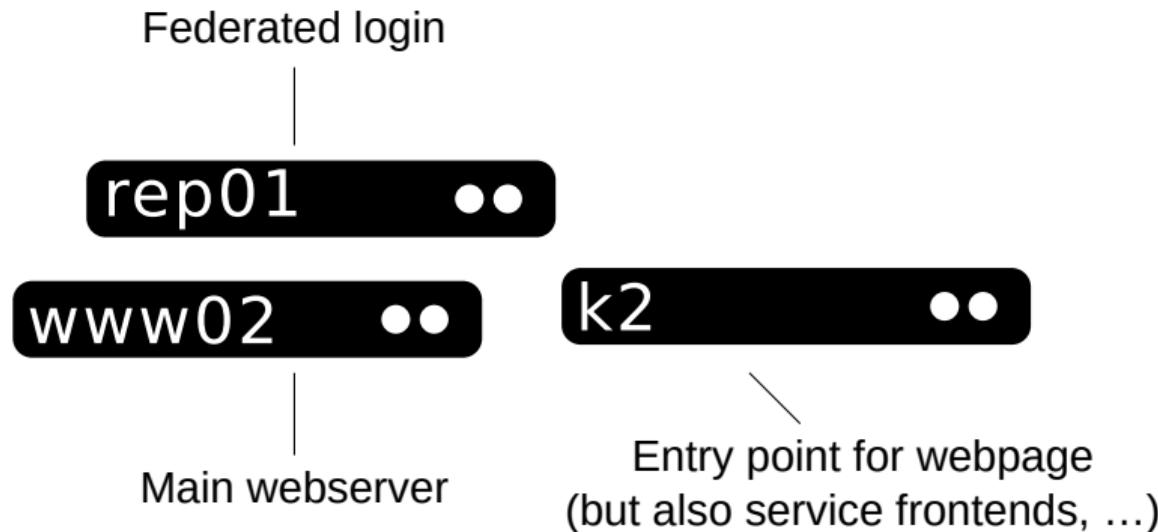


Entry point for webpage
(but also service frontends, ...)

Servers – Web



Servers – Web



Servers – Core Services

spraakbanken01p ●●

spraakbanken04p ●●

spraakbanken02p ●●

spraakbanken05p ●●

spraakbanken03p ●●

korpsrv ●●

Servers – Core Services

Strix backend



spraakbanken01p ●●

spraakbanken04p ●●

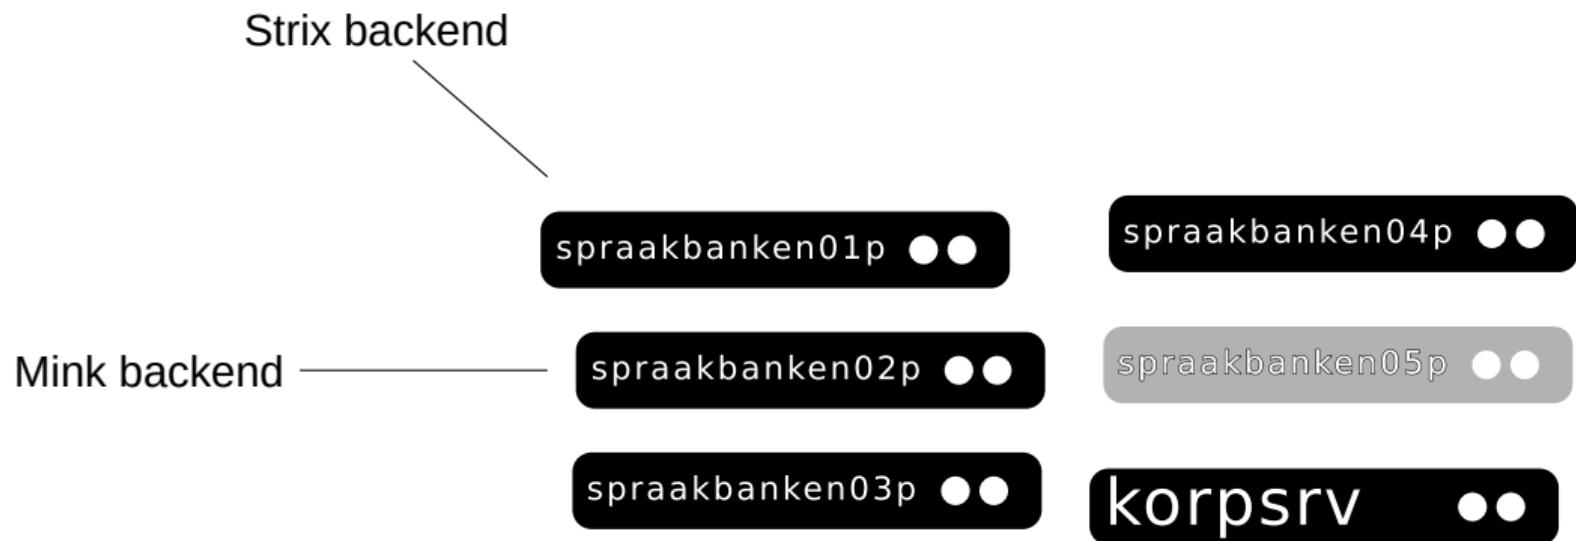
spraakbanken02p ●●

spraakbanken05p ●●

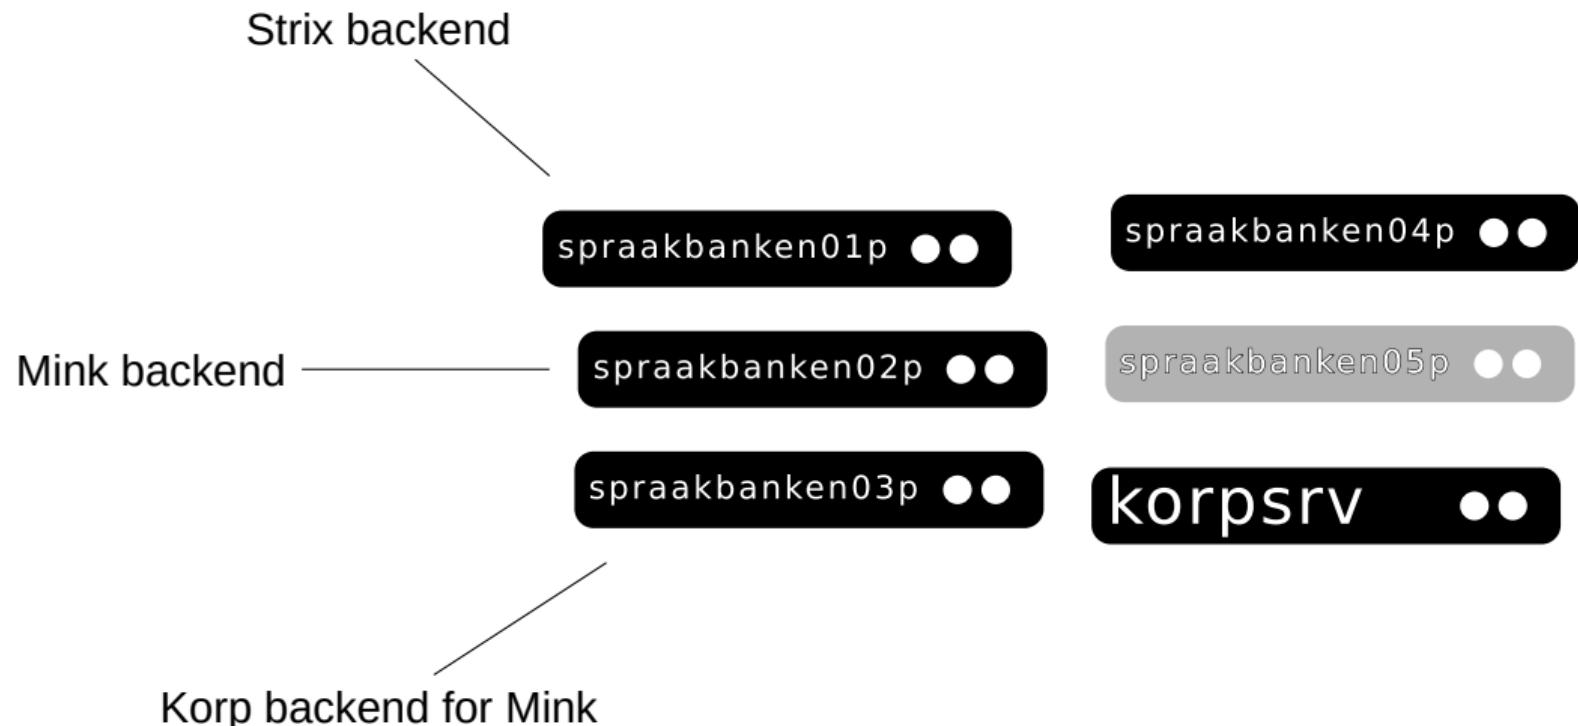
spraakbanken03p ●●

korpsrv ●●

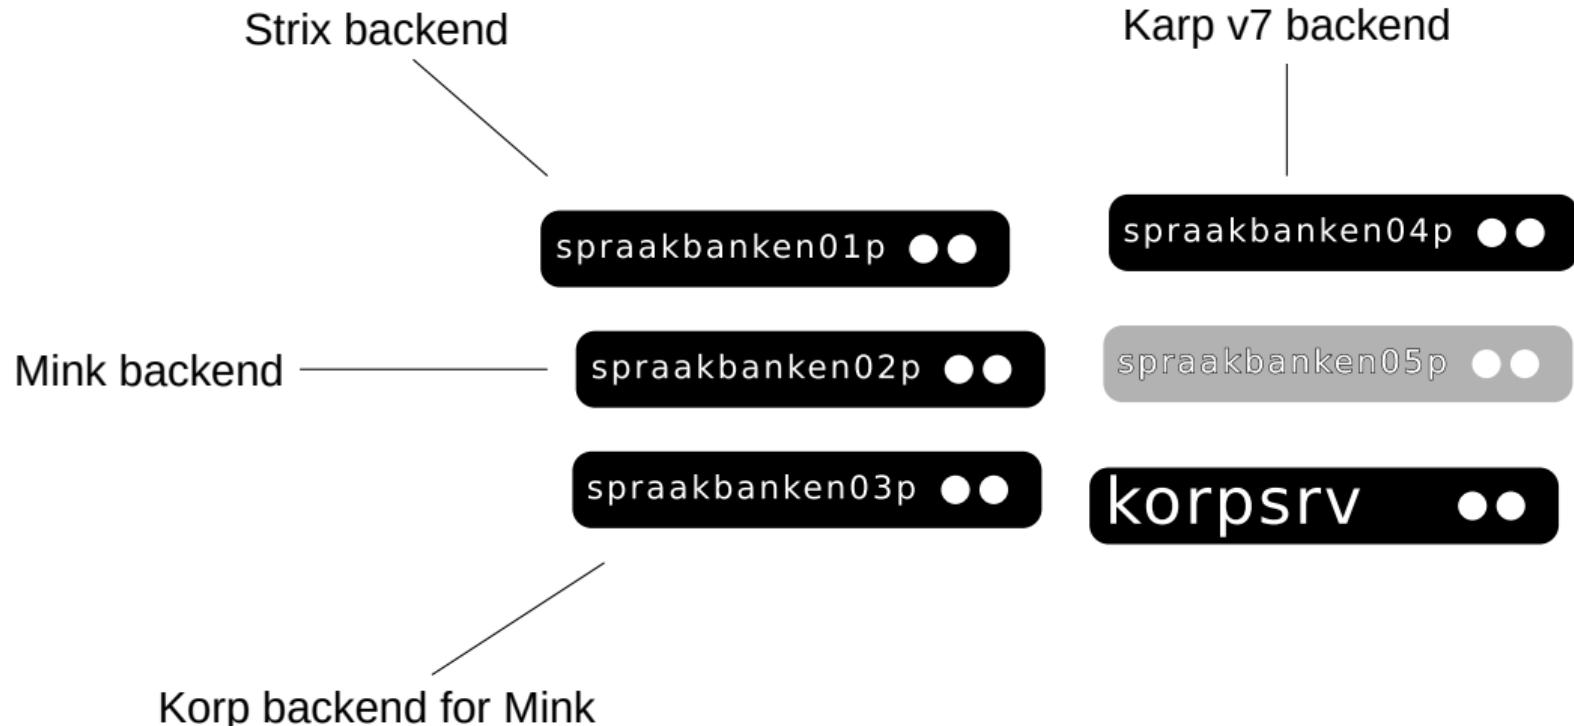
Servers – Core Services



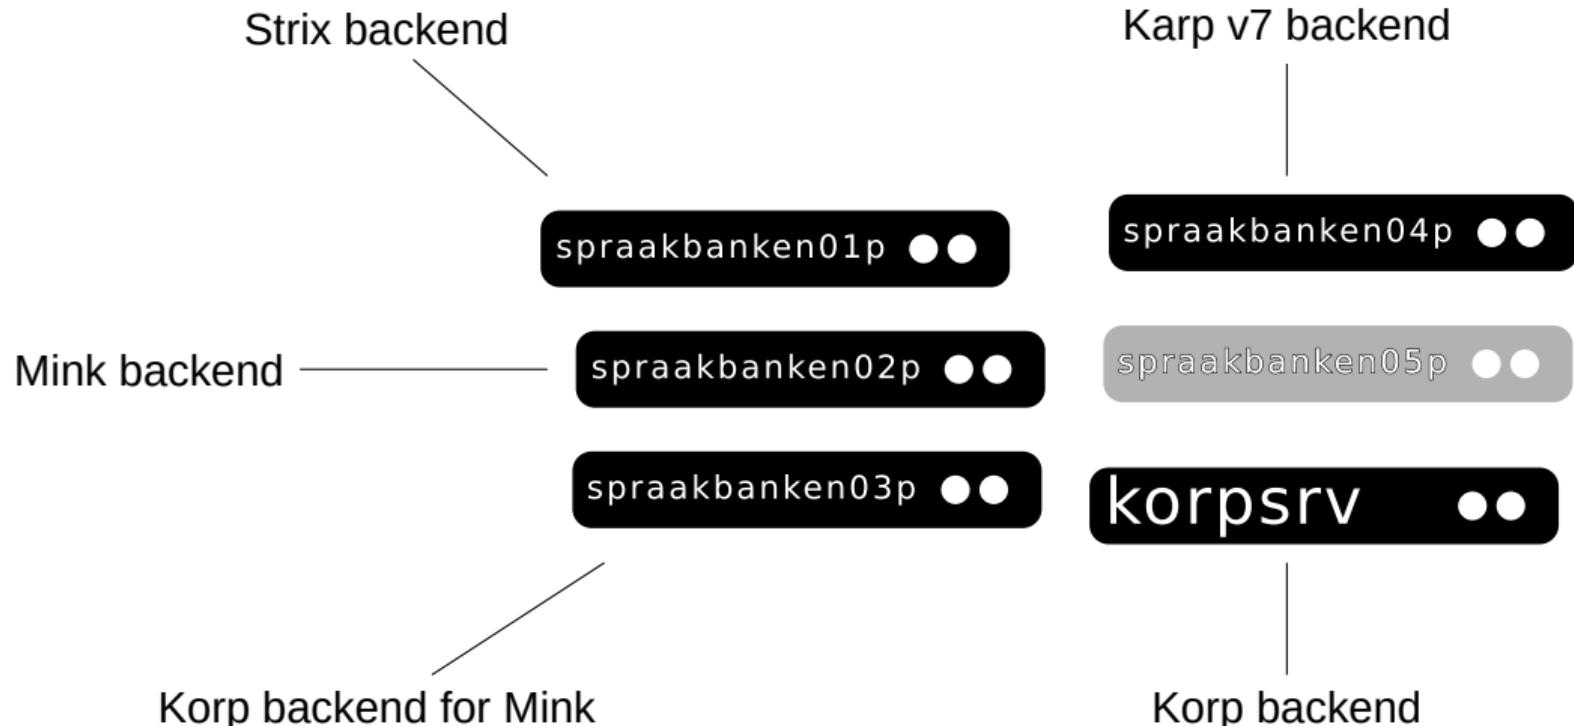
Servers – Core Services



Servers – Core Services



Servers – Core Services



Servers – Compute

kark ••

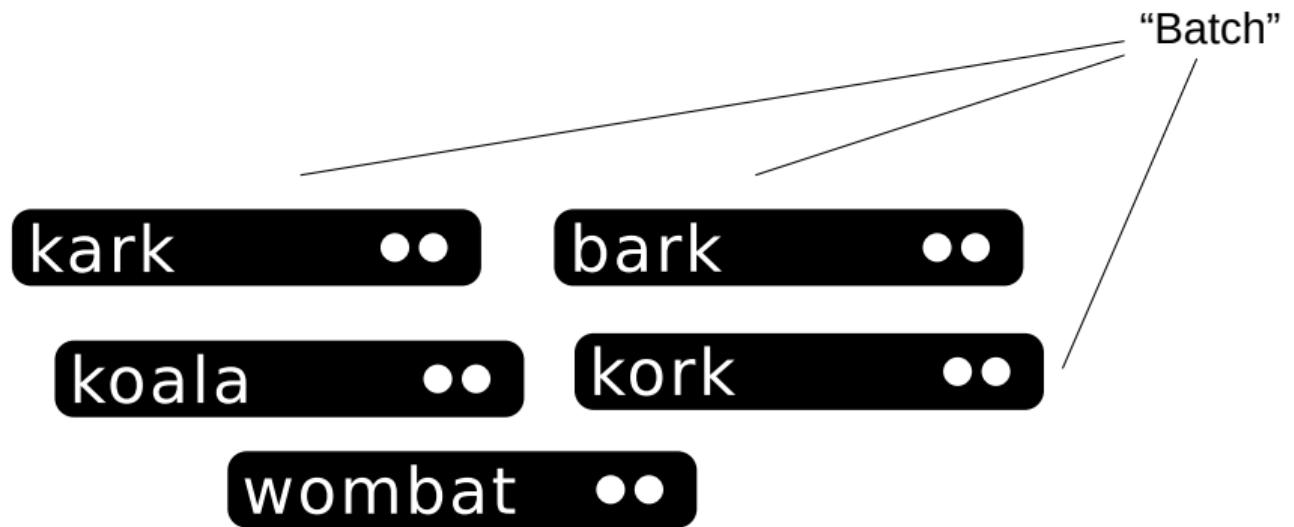
bark ••

koala ••

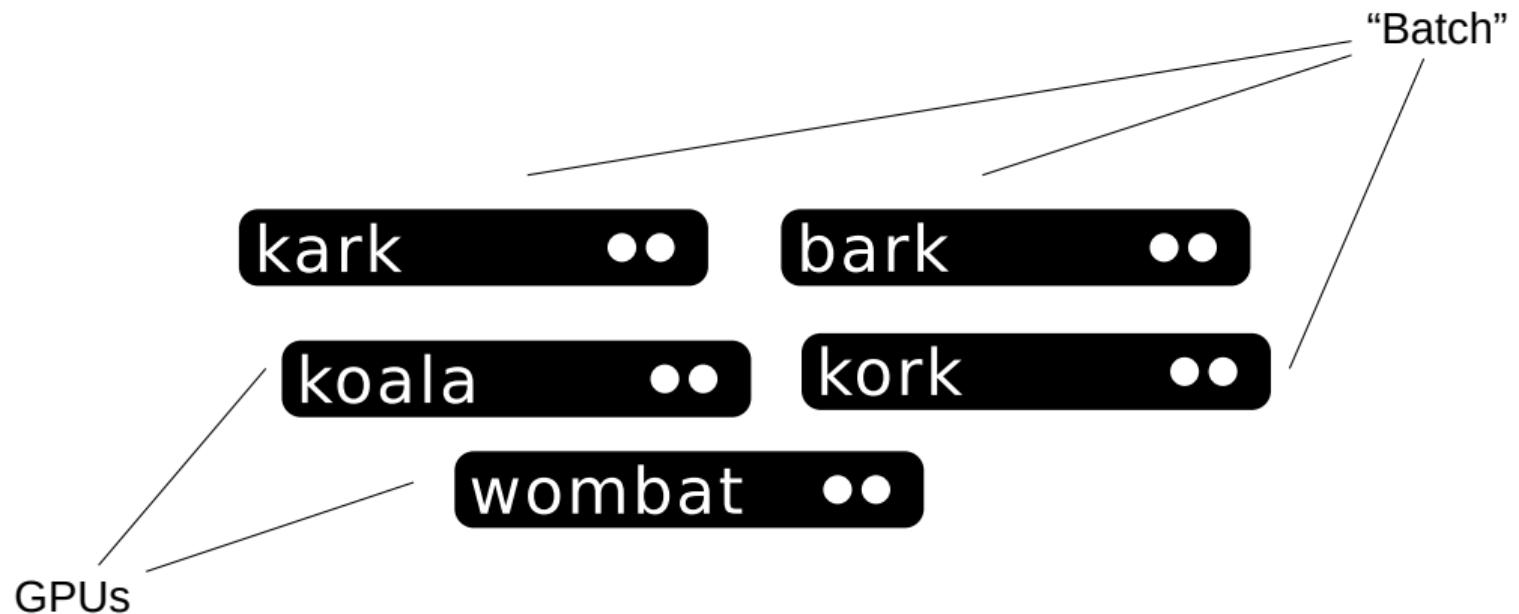
kork ••

wombat ••

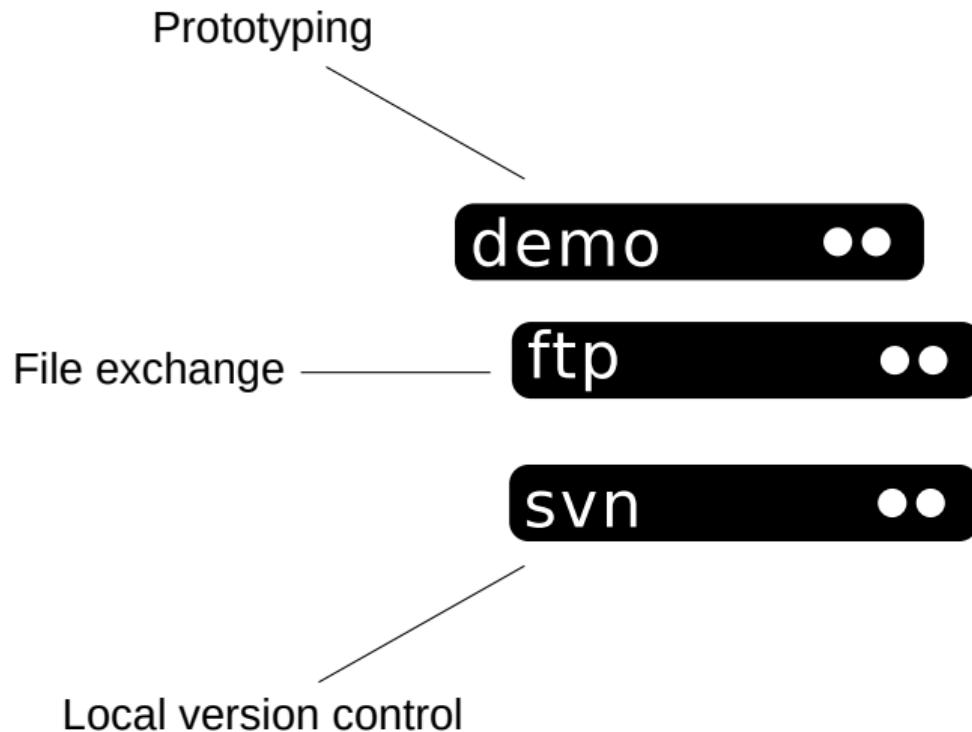
Servers – Compute



Servers – Compute



Servers – Development/Data exchange



Rack



More information

<https://spraakbanken.gu.se/om/internt/teknik/servrar>



Best practice

Just like when you are a guest somewhere

- Be nice
- Don't litter

Be nice

Don't occupy shared resources

- ✖ Storage is a shared resource, writing a lot of data to your home folder (or anywhere else on a network share) will slow down everything for everyone else ⇒ Store temporary data on local disk instead
- ✖ If start long-running processes be *nice*¹ to everyone else. Use *htop* to see what is going on on the server
- ✖ GPUs are very limited, consider checking if anyone else is using them before starting a new experiment (*nvidia-smi*)

¹ nice - run a program with modified scheduling priority

Don't litter

- ▶ Don't store everything forever (e.g. *virtualenvs*)
- ▶ Check and cleanup your (old) data (e.g. *old_home* on bark and kork, *home* on k2)
- ▶ Exclude data from backup (create *.nobackup* or *NOBACKUP* files)

Backups

We don't have a backup policy (yet)

:(
:(

Backups

We don't have a backup policy (yet)

:(
:(

⇒ We do not/cannot guarantee for the safety of
your data ATM

Backups

We don't have a backup policy (yet)

:(
:(

⇒ We do not/cannot guarantee for the safety of
your data ATM

But we do our best and backup all
the data we can

And we have RAID ;)

Backup storage



The Future?

At GU-IT:

- 100TB object storage
- compute environment in the cloud (and on premise) soon(?)

And here?

Questions?

Hands on

https://github.com/spraakbanken/work_meeting_autum_2024



Running a Whisper model on wombat

Host a container on demo

Thanks for your attention